

Introduction to Statistical Programming

Timothy Wong¹
James Gammerman²

March 19, 2018

¹`timothy.wong@hotmail.co.uk`

²`person@organisation.com`

Outline

- Opening

- Introduction to Data Analysis

 - Data Transformation

- Regression Models

 - Linear Regression

 - Poisson Regression

 - Logistic Regression

- Tree-based Methods

 - Decision Tree

 - Random Forest

What is the R language?

- ▶ Offers modern and sophisticated statistical algorithms
- ▶ Used by millions of analysts and researchers worldwide
- ▶ Has a thriving open-source community
- ▶ Enables big Data analytics

Easy to Use

- ▶ PROC REG = `lm()` or `glm()`
- ▶ PROC SQL = `%>%`
- ▶ PROC SORT = `arrange()`
- ▶ PROC MEANS = `mean()`, `sd()`
- ▶ PROC GPLOT = `plot()`, `ggplot()`, `autoplot()`

RStudio Server Pro

- RStudio is your integrated development environment (IDE)

The screenshot displays the RStudio Server Pro interface. The main editor window shows an R script with the following code:

```
722  
723 # Visualize the non-priority appointment roll on a chart  
724 getHead(model.appt.priority,  
725 seq.Date(as.Date("2018-07-01") - days(MAX_DAYS_AHEAD_PRIORITY), as.Date("2018-07-01") + days(20), 1),  
726 c("11P11", "11P12", "11P13", "11P14"),  
727 1), events) %>%  
728 ggplot(aes(x=days_ahead, y=num_appt_Pred, colour=operating_district_code)) + geom_line() + facet_wrap(~created_date, ncol = 7) +  
729 #scale_y_log10() +  
730 xlab("Number of rolled days") +  
731 ylab("Number of expected appointments for each initial contact")  
732  
733 # Calculate the number of appointments landed on each date.  
734 wmis.appt.all.xdf %>%  
735 group_by(operating_district_code, is_priority, appt_date) %>%  
736 summarise(total_actual_appt = sum(num_appt)) %>%  
737 rxxdfToDataFrame() %>%  
738 as.tibble() %>%  
739 readr::write_csv("wmis_appt_total.csv")  
740  
741 wmis.appt.total <- readr::read_csv("wmis_appt_total.csv")  
742  
743 # Plot a chart  
744 fitted.contact %>%  
745 filter(date >= seq.Date(from = TRAIN_START_DATE,  
746 to = TEST_END_DATE,  
747 by = 1)) %>%  
748 select(opernum_day_of_rollcode, wmo_code, date,  
749 contact_non_priority, contact_priority,  
750 fitted_priority_fit, fitted_priorroll,  
751 fitted_non_priority_fit, fitted_non_priority_sp_fit) %>%  
752  
753  
754  
755  
756  
757  
758  
759  
760  
761  
762  
763  
764  
765  
766  
767  
768  
769  
770  
771  
772  
773  
774  
775  
776  
777  
778  
779  
780  
781  
782  
783  
784  
785  
786  
787  
788  
789  
790  
791  
792  
793  
794  
795  
796  
797  
798  
799  
800  
801  
802  
803  
804  
805  
806  
807  
808  
809  
810  
811  
812  
813  
814  
815  
816  
817  
818  
819  
820  
821  
822  
823  
824  
825  
826  
827  
828  
829  
830  
831  
832  
833  
834  
835  
836  
837  
838  
839  
840  
841  
842  
843  
844  
845  
846  
847  
848  
849  
850  
851  
852  
853  
854  
855  
856  
857  
858  
859  
860  
861  
862  
863  
864  
865  
866  
867  
868  
869  
870  
871  
872  
873  
874  
875  
876  
877  
878  
879  
880  
881  
882  
883  
884  
885  
886  
887  
888  
889  
890  
891  
892  
893  
894  
895  
896  
897  
898  
899  
900  
901  
902  
903  
904  
905  
906  
907  
908  
909  
910  
911  
912  
913  
914  
915  
916  
917  
918  
919  
920  
921  
922  
923  
924  
925  
926  
927  
928  
929  
930  
931  
932  
933  
934  
935  
936  
937  
938  
939  
940  
941  
942  
943  
944  
945  
946  
947  
948  
949  
950  
951  
952  
953  
954  
955  
956  
957  
958  
959  
960  
961  
962  
963  
964  
965  
966  
967  
968  
969  
970  
971  
972  
973  
974  
975  
976  
977  
978  
979  
980  
981  
982  
983  
984  
985  
986  
987  
988  
989  
990  
991  
992  
993  
994  
995  
996  
997  
998  
999  
1000
```

The console window shows the output of the script, including the creation of the `wmis.appt.total` tibble and the first few rows of data:

```
appt_date = col_date(format = "%Y-%m-%d"),  
total_actual_appt = col_integer()  
> wmis.appt.total  
# A tibble: 354,604 x 4  
operating_district_code is_priority appt_date total_actual_appt  
  <chr> <lgl> <dttm> <int>  
1 11P11 FALSE 2015-07-01 0  
2 11P42 FALSE 2015-07-01 5  
3 65V04 FALSE 2015-07-01 0  
4 21Q75 FALSE 2015-07-01 2  
5 32R32 FALSE 2015-07-01 3  
6 11P42 FALSE 2015-07-01 5  
7 21Q74 FALSE 2015-07-01 3  
8 65V73 FALSE 2015-07-01 1  
9 70463 FALSE 2015-07-01 1  
10 70442 FALSE 2015-07-01 5  
# ... with 354,594 more rows
```

The Environment pane on the right shows the loaded objects:

- `wmis.appt.total`: 354604 obs. of 4 variables
- `wmis.contact`: 348196 obs. of 4 variables
- `wmis.ref.geogra`: 18318 obs. of 17 variables
- `wmis.wmo.mapping`: 173 obs. of 2 variables
- `wx.actual.daily`: 49648 obs. of 9 variables
- `wx.actual.hourly`: 1108932 obs. of 10 variables
- `wx.seasonal.non`: 219688 obs. of 7 variables
- `wx.seasonal.non`: 9158 obs. of 9 variables

The Files pane on the right shows the project files and folders, including `model.RData`, `model.Rproj`, `model.Rmd`, `model.Lappt_priority.rds`, `patch_map.Rmd`, `Sectors.xml`, `wmis_appt.all.xdf`, `wmis_appt_nonpriority.xdf`, and `wmis_appt_priority.xdf`.

Packages

- ▶ **CRAN** is the Comprehensive R Archive Network.
- ▶ User-contributed packages: source code, binaries, documentation.

```
# Install a new package with all its dependencies  
install.packages('ggplot2', dependencies = TRUE)  
# Load an installed package  
# (both lines are identical)  
library(ggplot2)  
library('ggplot2')
```

Packages

- ▶ **CRAN Task View** is a curated list of packages.
- ▶ <https://cran.r-project.org/web/views/>

```
library(ctv)
# Install a CRAN Task View
install.views("Econometrics")
# Update a CRAN Task View
update.views("Econometrics")
```

Vectors

- ▶ R is a vectorised programming language.
- ▶ Vector contains objects of the same data type.

```
# Create a vector of integers one to ten
```

```
myVec1 <- 1:10
```

```
# Find out the length of vector
```

```
length(myVec1)
```

```
## [1] 10
```

```
# Reverse the vector
```

```
rev(myVec1)
```

```
## [1] 10 9 8 7 6 5 4 3 2 1
```

```
# Create a custom vector of 10, 15, 20, 25, 30
```

```
myVec2 <- c(10, 15, 20, 25, 30)
```

```
myVec2
```

```
## [1] 10 15 20 25 30
```

```
# Create a vector of sequential numbers
```

```
myVec3 <- seq(from = -10, to = 10, by = 0.5)
```

```
myVec3
```

```
## [1] -10.0 -9.5 -9.0 -8.5 -8.0 -7.5 -7.0 -6.5 -6.0 -5.5 -5.0
```

```
## [12] -4.5 -4.0 -3.5 -3.0 -2.5 -2.0 -1.5 -1.0 -0.5 0.0 0.5
```

```
## [23] 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0 5.5 6.0
```

```
## [34] 6.5 7.0 7.5 8.0 8.5 9.0 9.5 10.0
```


Subsetting a Vector

- ▶ You can subset members from a vector.

```
# Select the second member of the vector  
myVec1[2]
```

```
## [1] 2
```

```
# Subset a range from the vector  
myVec1[2:4]
```

```
## [1] 2 3 4
```

```
# Subset specified elements  
myVec1[c(4,2,3)]
```

```
## [1] 4 2 3
```

Vectorised Operations

- Operations in R are vectorised.

```
# Arithmetic operations
```

```
myVec1 + 10
```

```
## [1] 11 12 13 14 15 16 17 18 19 20
```

```
myVec1 - 10
```

```
## [1] -9 -8 -7 -6 -5 -4 -3 -2 -1 0
```

```
myVec1 * 2
```

```
## [1] 2 4 6 8 10 12 14 16 18 20
```

```
myVec1 / 2
```

```
## [1] 0.5 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0
```

```
myVec1 ^ 2
```

```
## [1] 1 4 9 16 25 36 49 64 81 100
```

```
log(myVec1)
```

```
## [1] 0.0000000 0.6931472 1.0986123 1.3862944 1.6094379 1.7917595 1.9459101
```

```
## [8] 2.0794415 2.1972246 2.3025851
```

Looping

- ▶ Looping can be slow.
- ▶ Always try to vectorise your code.

```
# Vectorised operation is fast
system.time({
  myResult <- 1:100000 * 2
})

##      user  system elapsed
##         0         0         0

# Looping is quite slow
system.time({
  myResult <- sapply(1:100000, function(x){ x * 2 })
})

##      user  system elapsed
##    0.05    0.02    0.06

# Appending to vector is much slower
system.time({
  myResult <- c()
  for(i in 1:100000){
    myResult <- c(myResult, i * 2)
  }
})

##      user  system elapsed
##   16.01    0.15    16.23
```

Functions

- Functions are vectorised.

```
# Defines a custom function
myFunc <- function(x) {
  x * 2
}
# Execute the function with one input
myFunc(5)

## [1] 10

# Execute the function with an integer vector
myFunc(1:10)

## [1] 2 4 6 8 10 12 14 16 18 20
```

Character Vectors

- Vector can also contain character objects.

```
# Vector can contain character objects
myVec4 <- c('Bill', 'Mark', 'Steve', 'Jeff', 'Larry')
myVec4

## [1] "Bill" "Mark" "Steve" "Jeff" "Larry"

# Constant character vectors in R
LETTERS

## [1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O" "P" "Q"
## [18] "R" "S" "T" "U" "V" "W" "X" "Y" "Z"

letters

## [1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q"
## [18] "r" "s" "t" "u" "v" "w" "x" "y" "z"

month.name

## [1] "January" "February" "March" "April" "May"
## [6] "June" "July" "August" "September" "October"
## [11] "November" "December"

month.abb

## [1] "Jan" "Feb" "Mar" "Apr" "May" "Jun" "Jul" "Aug" "Sep" "Oct" "Nov"
## [12] "Dec"
```

Vectors - Other Data Types

- Vector can contain objects of any data type.

```
# This is a vector of character objects
myVec5 <- c('2017-07-13',
            '2017-10-11',
            '2017-11-21',
            '2018-01-16',
            '2018-03-27')

# This is a vector of date objects
myVec6 <- as.Date(myVec5)
# Compute the day of week - returns a vector of characters
# Notice that these are all vectorised functions
weekdays(myVec6)

## [1] "Thursday" "Wednesday" "Tuesday"    "Tuesday"    "Tuesday"
```

List

- List is a generic container for objects of different data types

```
myFavBook <- list(title = 'R for Data Science',
  authors = c('Garrett Grolemund', 'Hadley Wickham'),
  publishDate = as.Date('2016-12-12'),
  price = 18.17,
  currency = 'USD',
  edition = 1,
  isbn = 1491910399)

myFavBook

## $title
## [1] "R for Data Science"
##
## $authors
## [1] "Garrett Grolemund" "Hadley Wickham"
##
## $publishDate
## [1] "2016-12-12"
##
## $price
## [1] 18.17
##
## $currency
## [1] "USD"
##
## $edition
## [1] 1
##
## $isbn
## [1] 1491910399
```

Subsetting a List

- ▶ You can subset a particular member from a list

```
# Select a named member of a list  
# Using the dollar sign, followed by name without bracket  
myFavBook$title  
  
## [1] "R for Data Science"  
  
# Using double squared brackets with member's name as string  
myFavBook[['authors']]  
  
## [1] "Garrett Grolemund" "Hadley Wickham"  
  
# Select the fourth member in the list  
myFavBook[[4]]  
  
## [1] 18.17
```


Special Numbers in R

```
# Pi is constant 3.14159...
pi

## [1] 3.141593

# One divided by zero is infinity
1/0

## [1] Inf

# Negative number divided by zero is negative infinity
-1/0

## [1] -Inf

# Infinity divided by infinity is Not-a-Number (NaN)
Inf/Inf

## [1] NaN

# Not available (NA) plus one is still NA
NA + 1

## [1] NA

# Effects of different special numbers
c(5, 10, 15, NA, 25, 30, NaN, 35, 40, Inf, 50, -Inf, 60) / 5

## [1] 1 2 3 NA 5 6 NaN 7 8 Inf 10 -Inf 12
```

Data Frame

- ▶ Table with rows (observations) and columns (variables).
- ▶ Analogous to an Excel workbook.

[illegible]

Data Frame

```
myFavMovies1
```

```
##           title year    box      bondActor
## 1           Dr. No 1962   59.5      Sean Connery
## 2         Goldfinger 1964  125.0      Sean Connery
## 3 Diamonds are Forever 1971 120.0      Sean Connery
## 4           Moonraker 1979 210.3      Roger Moore
## 5 The Living Daylights 1987 191.2 Timothy Dalton
## 6           GoldenEye 1995 355.0 Pierce Brosnan
## 7       Casino Royale 2006 599.0   Daniel Craig
```

Tibble

- ▶ Similar to traditional data frame.
- ▶ tibble is the modern standard in R.

```
library(dplyr)
myFavMovies2 <- tibble(title = c('Dr. No',
                                'Goldfinger',
                                'Diamonds are Forever',
                                'Moonraker',
                                'The Living Daylights',
                                'GoldenEye',
                                'Casino Royale'),
  year = c(1962, 1964, 1971, 1979,
           1987, 1995, 2006),
  box = c(59.5, 125, 120, 210.3,
          191.2, 355, 599),
  bondActor = c('Sean Connery',
                'Sean Connery',
                'Sean Connery',
                'Roger Moore',
                'Timothy Dalton',
                'Pierce Brosnan',
                'Daniel Craig'))
```

Tibble

```
myFavMovies2
```

```
## # A tibble: 7 x 4
```

##		title	year	box	bondActor
##		<chr>	<dbl>	<dbl>	<chr>
## 1		Dr. No	1962	59.5	Sean Connery
## 2		Goldfinger	1964	125.0	Sean Connery
## 3	Diamonds are Forever		1971	120.0	Sean Connery
## 4		Moonraker	1979	210.3	Roger Moore
## 5	The Living Daylights		1987	191.2	Timothy Dalton
## 6		GoldenEye	1995	355.0	Pierce Brosnan
## 7		Casino Royale	2006	599.0	Daniel Craig

Subsetting a Tibble

```
# Get one column by name  
myFavMovies2[['title']]  
myFavMovies2$title  
# Get a range of columns by position ID  
myFavMovies2[, 1:2]  
myFavMovies2[1:2]  
# Get rows 1 to 3  
myFavMovies2[1:3, ]  
# Get the 'year' variable of row 1-3  
myFavMovies2[1:3, 'year']  
# Get the 'title' and 'year' variables of row 4-7  
myFavMovies2[4:7, c('title', 'year')]
```

Simple Regression

Univariate linear regression model

$$\hat{y}_i = \beta_0 + \beta_1 x_i$$

- ▶ Analogous to a straight line $y = mx + c$
- ▶ Can be chained with M dependent variables (Multivariate)

$$\hat{y}_i = \beta_0 + \sum_{m=1}^M \beta_m x_{m,i}$$

- ▶ Residual term $\epsilon_i = y_i - \hat{y}_i$ assumed to be Gaussian

$$\epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

- ▶ Also known as ordinary least squared (OLS) regression

Linear Regression in R

```
# Build a univariate linear model
# These two lines are equivalent
myModel1 <- lm(mpg ~ wt, mtcars)
myModel1 <- lm(formula = mpg ~ wt, data = mtcars)
# Read the model summary
summary(myModel1)

##
## Call:
## lm(formula = mpg ~ wt, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.5432 -2.3647 -0.1252  1.4096  6.8727
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   37.2851     1.8776  19.858 < 2e-16 ***
## wt           -5.3445     0.5591  -9.559 1.29e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.046 on 30 degrees of freedom
## Multiple R-squared:  0.7528, Adjusted R-squared:  0.7446
## F-statistic: 91.38 on 1 and 30 DF,  p-value: 1.294e-10
```


More Linear Regression Models

Multivariate linear model Additional independent variables can be chained using the + symbol. Categorical variables can be encoded as dummy on-the-fly using `factor()`.

```
myModel2 <- lm(mpg ~ wt + hp + qsec + factor(am), mtcars)
```

Polynomial term Model can become more flexible when an independent variable is converted into polynomial terms. Use the `poly()` function.

```
myModel3 <- lm(mpg ~ wt + qsec + factor(am) +  
               poly(hp, 3), mtcars)
```

Interaction term Two variables can be combined to create synergy effect. The * symbol is used to combine variables together.

```
myModel4 <- lm(mpg ~ wt * hp + qsec + factor(am), mtcars)
```

Regression Diagnostics

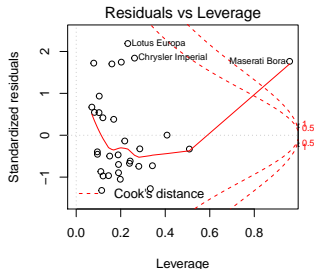
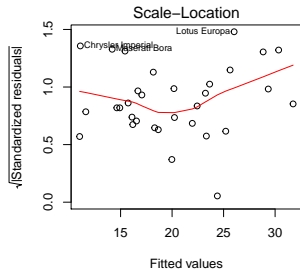
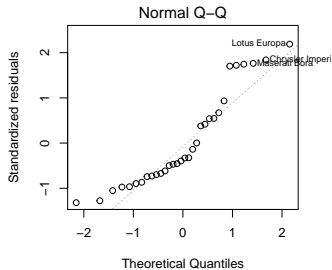
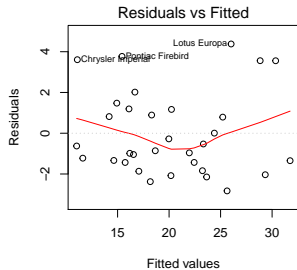
Residuals vs Fitted Checks for non-linear relationship. Look for a near horizontal line.

Normal Quantile-Quantile It aligns model residuals against a theoretical normal distribution. If the residuals spread along a straight diagonal line on the Q-Q plot, it suggests that the residuals are normally distributed.

Scale-Location Checks for homoscedasticity and heteroscedasticity. It is homoscedastic if observations scatter without any observable pattern.

Residual vs Leverage (Cook's Distance) Identifies observations having strong influence to the model.

Diagnostics Plots

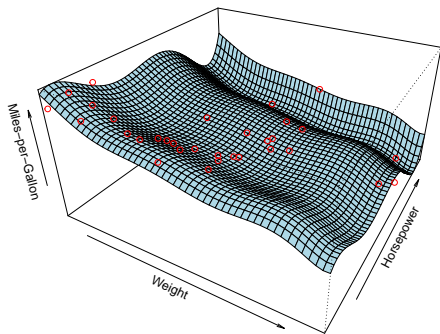


Overfitting

- ▶ Flexible models are prone to overfitting.
- ▶ Overfitting makes the model less generalisable.
- ▶ Solution
 - ▶ Use less flexible methods.
 - ▶ Impose regularisation.


Overfitting: Visual Explanation

$$\hat{y} = \beta_0 + \sum_{j=1}^8 \beta_{wtj} x_{wt}^j + \sum_{k=1}^5 \beta_{hp_k} x_{hp}^k$$



Poisson Distribution

- ▶ Count of distinct events are drawn from Poisson distribution.
 - ▶ Always positive.
 - ▶ In most cases they are integers.
- ▶ e.g. Number of people in a room, number of flights delayed per day... etc.




Testing for Poisson Distribution

- Chi-square goodness of fit test.

```
# Performs the Chi-squared goodness-of-fit test.  
# It checks whether the variable is drawn from a Poisson distribution.  
library(vcd)  
gf <- goodfit(mtcars$carb, type= "poisson", method= "ML")  
# Checks the statistical p-value of the goodness-of-fit test.  
# If  $p \leq 0.05$  then it is safe to say that the variable is Poisson.  
summary(gf)  
  
##  
##      Goodness-of-fit test for poisson distribution  
##  
##              X^2 df      P(> X^2)  
## Likelihood Ratio 20.53973  4 0.0003906369
```

Goodness of Fit Plot

```
# Plots the observed frequency vs theoretical Poisson distribution.  
# The hanging bars should fill the space if it was perfectly Poisson.  
plot(gf)
```



Poisson Regression


```
# Build a Poisson model to predict the number of carburetors in a car.
myPoissonModel <- glm(carb ~ hp + wt + factor(am),
                      family="poisson",
                      data=mtcars)

# Read the model summary
summary(myPoissonModel)

##
## Call:
## glm(formula = carb ~ hp + wt + factor(am), family = "poisson",
##      data = mtcars)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.91420  -0.48423  -0.07246   0.19252   1.26155
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.418081   0.604211  -0.692   0.4890
## hp           0.004316   0.001880   2.296   0.0217 *
## wt           0.179583   0.191352   0.938   0.3480
## factor(am)1  0.393750   0.324978   1.212   0.2257
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 27.043  on 31  degrees of freedom
## Residual deviance: 10.798  on 28  degrees of freedom
## AIC: 108.16
##
## Number of Fisher Scoring iterations: 4
```

Recursive Partitioning

- ▶ Cut off point is denoted as s .
- ▶ Divides data into regions (leaves) $\mathcal{R}_1, \mathcal{R}_2, \mathcal{R}_3, \dots$ recursively.
- ▶ Works with real values as well as categorical variables.
- ▶ Large tree risks overfitting
 - ▶ Removes weaker leaves.
 - ▶ Regularisation.



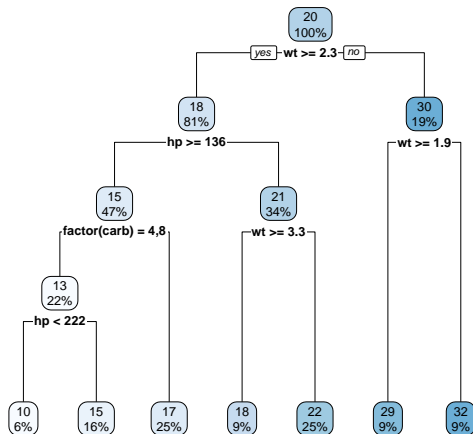
Decision Tree

- Trees can be trained with a formula and optional control parameters.

```
# Load the rpart package for recursive partitioning  
library(rpart)  
# Build a decision tree to predict mpg  
myTree <- rpart(formula = mpg ~ wt + hp +  
                 factor(carb) +  
                 factor(am),  
                 data = mtcars,  
                 control = rpart.control(minsplit=5))  
# Read the detailed summary of the tree  
summary(myTree)
```

Decision Tree: Visualisation

```
# Load the rpart.plot package for tree visualisation  
library(rpart.plot)  
rpart.plot(myTree)
```




Tree Pruning

```
printcp(myTree)

##
## Regression tree:
## rpart(formula = mpg ~ wt + hp + factor(carb) + factor(am), data = mtcars,
##       control = rpart.control(minsplit = 5))
##
## Variables actually used in tree construction:
## [1] factor(carb) hp      wt
##
## Root node error: 1126/32 = 35.189
##
## n= 32
##
##      CP nsplit rel error  xerror   xstd
## 1 0.652661    0 1.000000 1.10529 0.265228
## 2 0.178618    1 0.347339 0.66000 0.169220
## 3 0.046269    2 0.168721 0.46702 0.118747
## 4 0.026109    3 0.122453 0.43399 0.118450
## 5 0.022593    4 0.096343 0.40681 0.086571
## 6 0.011989    5 0.073751 0.40145 0.089376
## 7 0.010000    6 0.061762 0.39935 0.089194

plotcp(myTree)
```



Random Forest

- ▶ Consists of many decision trees
 - ▶ Randomly selected variables will be used in each split
 - ▶ Usually no need to prune them (all trees are allowed to grow big)
- ▶ M trees in a random forest produces M predictions
 - ▶ Final prediction is calculated as mean value for regression problem
 - ▶ Classification problem will use most the common label (majority voting)

Training a Random Forest

```
library(randomForest)
library(dplyr)
# Build a random forest with 1000 trees
# Each tree has 2 randomly selected variables
# You can change the parameters
myForest <- randomForest(mpg ~ wt + hp + carb + am,
                          ntree = 1000,
                          mtry = 2,
                          data = mtcars %>% mutate(carb = factor(carb),
                                                    am = factor(am)))

# Plot the error as the forest expands
plot(myForest)
```

