

Introduction to Statistical Programming

Timothy Wong¹
James Gammerman²

August 22, 2018

¹`timothy.wong@hotmail.co.uk`

²`jgammerman@gmail.com`

Outline

- R Ecosystem

- Data Transformation

- Regression Models

 - Linear Regression

 - Poisson Regression

 - Logistic Regression

- Tree-based Methods

 - Decision Tree

 - Random Forest

- Neural Networks

 - Multilayer Perceptron

- Time Series Analysis

 - Auto-Correlation Function

 - Decomposition

 - ARIMA Model

- Survival Analysis

 - Kaplan-Meier Estimator

 - Cox Proportional Hazard Model

- Unsupervised Learning

 - K-means Clustering

 - Hierarchical Clustering

R Ecosystem

What is the R language?

- ▶ Offers modern and sophisticated statistical algorithms
- ▶ Used by millions of analysts and researchers worldwide
- ▶ Has a thriving open-source community
- ▶ Enables big data analytics

Easy to Use

- ▶ PROC REG = `lm()` or `glm()`
- ▶ PROC SQL = `%>%`
- ▶ PROC SORT = `arrange()`
- ▶ PROC MEANS = `mean()`, `sd()`
- ▶ PROC GPLOT = `plot()`, `ggplot()`, `autoplot()`

RStudio Server Pro

- RStudio is your integrated development environment (IDE)

The screenshot displays the RStudio Server Pro interface. The main editor window shows an R script with the following code:

```
722 # Visualize the non-priority appointment roll on a chart
723
724 getHead(model.appt.priority,
725         seq.Date(as.Date("2018-07-01") - days(MAX_DAYS_AHEAD_PRIORITY), as.Date("2018-07-01") + days(20), 1),
726         c("11P11", "11P12", "11P13", "11P14"),
727         14, events) %>%
728 ggplot(aes(x=days_ahead, y=num_appt_pred, colour=operating_district_code)) + geom_line() + facet_wrap(~created_date, ncol = 7) +
729     scale_y_log10() +
730     xlab("Number of rolled days") +
731     ylab("Number of expected appointments for each initial contact")
732
733 # Calculate the number of appointments landed on each date.
734
735 wmis.appt.all.xdf %>%
736   group_by(operating_district_code, is_priority, appt_date) %>%
737   summarise(total_actual_appt = sum(num_appt)) %>%
738   rxxdfToDataFrame() %>%
739   as.tibble() %>%
740   readr::write_csv("wmis_appt_total.csv")
741
742 wmis.appt.total <- readr::read_csv("wmis_appt_total.csv")
743
744 # Plot a chart
745 fitted.contact %>%
746   filter(date >= seq.Date(from = TRAIN_START_DATE,
747                           to = TEST_END_DATE,
748                           by = 1)) %>%
749   select(opernum_day_of_rollcode, wmo_code, date,
750          contact_non_priority, contact_priority,
751          fitted_priority_fit, fitted_priorroll)
752
753 fitted_non_priority_fit, fitted_non_priority.sp.fit) %>%
```

The console output shows the execution of the script, including the creation of the `wmis.appt.total` tibble and the resulting data structure:

```
appt_date = col_date(format = "%Y-%m-%d"),
total_actual_appt = col_integer()
> wmis.appt.total
# A tibble: 354,604 x 4
  operating_district_code is_priority appt_date total_actual_appt
  <chr> <lgl> <dttm> <int>
1 11P11 FALSE 2015-07-01 0
2 11P42 FALSE 2015-07-01 5
3 65V04 FALSE 2015-07-01 0
4 21Q75 FALSE 2015-07-01 2
5 32R32 FALSE 2015-07-01 3
6 11P42 FALSE 2015-07-01 5
7 21Q74 FALSE 2015-07-01 3
8 65V73 FALSE 2015-07-01 1
9 70463 FALSE 2015-07-01 1
10 70442 FALSE 2015-07-01 5
# ... with 354,594 more rows
```

The Environment pane on the right shows the loaded objects and their sizes:

- Global Environment
- wmis.appt.total: 354604 obs. of 4 variables
- wmis.contact: 348196 obs. of 4 variables
- wmis.ref.geogra_: 18318 obs. of 17 variables
- wmis.wmo.mapping: 173 obs. of 2 variables
- wx.actual.daily: 49648 obs. of 9 variables
- wx.actual.hourly: 1108932 obs. of 10 variables
- wx.seasonal.non_: 219688 obs. of 7 variables
- wx.seasonal.non_: 9158 obs. of 9 variables

The Files pane on the right shows the project files and folders, including `gitignore`, `History`, `efods_actual_hourly.csv`, `events.csv`, `gas-repair-demand.Rproj`, `gr.appt.hol`, `gr.contact.hol`, `mapping_l6.csv`, `model.RData`, `model.Rmd`, `model.Lappt.priority.rds`, `patch_map.Rmd`, `Sectors.xml`, `wmis_appt.all.xdf`, `wmis_appt_nonpriority.xdf`, and `wmis_appt_priority.xdf`.

Packages

- ▶ **CRAN** is the Comprehensive R Archive Network.
- ▶ User-contributed packages: source code, binaries, documentation.

```
# Return all installed packages  
installed.packages()  
# Install a new package and all its dependencies from CRAN  
# This will install to the default library location  
install.packages("ggplot2", dependencies = TRUE)  
# Load an installed package  
# (both lines are identical)  
library(ggplot2)  
library("ggplot2")
```

Packages

- ▶ **CRAN Task View** is a curated list of packages.
- ▶ Useful guide to get started with R.
- ▶ <https://cran.r-project.org/web/views/>

Variable Assignment

- ▶ Assign variables using the <- symbol.
- ▶ Do not use reserved words. This will confuse the interpreter.

```
# Assign variables  
myVarX <- 5  
myVarY <- 20  
# Perform multiplication  
myVarX * myVarY  
# Look at the reserved words  
?Reserved
```

Vectors

- ▶ R is a vectorised programming language.
- ▶ Vector contains elements of the same data type.

```
myVec1 <- 1:10
# Find out the length of vector
length(myVec1)

## [1] 10

# Reverse the vector
# This does not change the value of myVec1
rev(myVec1)

## [1] 10 9 8 7 6 5 4 3 2 1

# Create a custom vector of 10, 15, 20, 25, 30
myVec2 <- c(10, 15, 20, 25, 30)
myVec2

## [1] 10 15 20 25 30

# Create a vector of sequential numbers with increment 0.5
myVec3 <- seq(from = -2, to = 2, by = 0.5)
myVec3

## [1] -2.0 -1.5 -1.0 -0.5 0.0 0.5 1.0 1.5 2.0

# Elements of a vector can be named
myVec4 <- c(`New York` = 8.5,
            `London` = 8.6,
            `Moscow` = 11.9)
myVec4

## New York    London    Moscow
##      8.5      8.6     11.9
```

Subsetting a Vector

- You can subset elements from a vector.

```
# Select the second element of the vector  
myVec2[2]
```

```
## [1] 15
```

```
# Subset a range from the vector  
myVec2[2:4]
```

```
## [1] 15 20 25
```

```
# Subset specified elements  
myVec2[c(4,2,3)]
```

```
## [1] 25 15 20
```

```
# Subset named element of a vector  
myVec4["New York"]
```

```
## New York
```

```
##      8.5
```

Vectorised Operations

- Operations in R are vectorised.

```
# Arithmetic operations
```

```
myVec1 + 10
```

```
## [1] 11 12 13 14 15 16 17 18 19 20
```

```
myVec1 - 10
```

```
## [1] -9 -8 -7 -6 -5 -4 -3 -2 -1 0
```

```
myVec1 * 2
```

```
## [1] 2 4 6 8 10 12 14 16 18 20
```

```
myVec1 / 2
```

```
## [1] 0.5 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0
```

```
myVec1 ^ 2
```

```
## [1] 1 4 9 16 25 36 49 64 81 100
```

```
log(myVec1)
```

```
## [1] 0.0000000 0.6931472 1.0986123 1.3862944 1.6094379 1.7917595 1.9459101
```

```
## [8] 2.0794415 2.1972246 2.3025851
```

Character Vectors

- Vector can also contain character objects.

```
# Vector can contain character objects
myVec4 <- c("Bill", "Mark", "Steve", "Jeff", "Larry")
myVec4

## [1] "Bill" "Mark" "Steve" "Jeff" "Larry"

# Constant character vectors in R
LETTERS

## [1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O" "P" "Q"
## [18] "R" "S" "T" "U" "V" "W" "X" "Y" "Z"

letters

## [1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q"
## [18] "r" "s" "t" "u" "v" "w" "x" "y" "z"

month.name

## [1] "January" "February" "March" "April" "May"
## [6] "June" "July" "August" "September" "October"
## [11] "November" "December"

month.abb

## [1] "Jan" "Feb" "Mar" "Apr" "May" "Jun" "Jul" "Aug" "Sep" "Oct" "Nov"
## [12] "Dec"
```

Date and Date/Time Vectors

- ▶ Date is a data type in base R.
- ▶ The package lubridate extends date/time functionalities.
- ▶ Enables easier date/time manipulation.

```
# This is a vector of Date objects
myVec5 <- as.Date(c("2017-07-13",
                    "2017-10-11",
                    "2017-11-21",
                    "2018-01-16",
                    "2018-03-27"))

# Load the lubridate package
# Use the function ymd_hms() to parse date/time with timezone
# Returns a vector of POSIXct (date/time) object
library(lubridate)
myVec6 <- ymd_hms(c("2017-07-13 09:30:00",
                    "2017-10-11 08:00:00",
                    "2017-11-21 10:00:00",
                    "2018-01-16 11:30:00",
                    "2018-03-27 12:00:00"),
                  tz = "Europe/London")

# Date/time manipulation applied to a vector
myVec7 <- myVec6 + hours(1) + minutes(30)
# Compute the day of week - returns a vector of characters
weekdays(myVec7)

## [1] "Thursday" "Wednesday" "Tuesday" "Tuesday" "Tuesday"
```

Logical Operators

- ▶ Apply logical operators on vector objects.

```
# Find all values greater than 5 - returns a vector of logical values  
myVec1 > 5
```

```
## [1] FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE
```

```
# Find all values equal to 7  
myVec1 == 7
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE
```

```
# Find all values matching 2,4,6 and 8  
myVec1 %in% c(2,4,6,8)
```

```
## [1] FALSE TRUE FALSE TRUE FALSE TRUE FALSE TRUE FALSE FALSE
```

```
# Find all values between 2 and 7  
myVec1 >= 2 & myVec1 <= 7
```

```
## [1] FALSE TRUE TRUE TRUE TRUE TRUE TRUE FALSE FALSE FALSE
```

```
# Find all values equal to 7 or equal to 8  
myVec1 == 7 | myVec1 == 8
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE TRUE TRUE FALSE FALSE
```

Special Numbers in R

```
# Pi is constant 3.14159...
pi

## [1] 3.141593

# One divided by zero is infinity
1/0

## [1] Inf

# Negative number divided by zero is negative infinity
-1/0

## [1] -Inf

# Infinity divided by infinity is Not-a-Number (NaN)
Inf/Inf

## [1] NaN

# Not available (NA) plus one is still NA
NA + 1

## [1] NA

# Effects of different special numbers
c(5, 10, 15, NA, 25, 30, NaN, 35, 40, Inf, 50, -Inf, 60) / 5

## [1] 1 2 3 NA 5 6 NaN 7 8 Inf 10 -Inf 12
```


List

- List is a generic container for objects of different data types

```
myFavBook <- list(title = "R for Data Science",
  authors = c("Garrett Grolemund", "Hadley Wickham"),
  publishDate = as.Date("2016-12-12"),
  price = 18.17,
  currency = "USD",
  edition = 1,
  isbn = 1491910399)

myFavBook

## $title
## [1] "R for Data Science"
##
## $authors
## [1] "Garrett Grolemund" "Hadley Wickham"
##
## $publishDate
## [1] "2016-12-12"
##
## $price
## [1] 18.17
##
## $currency
## [1] "USD"
##
## $edition
## [1] 1
##
## $isbn
## [1] 1491910399
```

Subsetting a List

- ▶ You can subset a particular member from a list

```
# Select a named element of a list  
# Use the dollar sign, followed by name without bracket  
myFavBook$title  
  
## [1] "R for Data Science"  
  
# Use double squared brackets with element's name as character  
myFavBook[["authors"]]  
  
## [1] "Garrett Grolemund" "Hadley Wickham"  
  
# Select the fourth element in the list  
myFavBook[[4]]  
  
## [1] 18.17
```

Data Frame

- ▶ Table with rows (observations) and columns (variables).
- ▶ Analogous to an Excel workbook.

[illegible]

Data Frame

```
myFavMovies1
```

```
##           title year    box      bondActor
## 1           Dr. No 1962   59.5      Sean Connery
## 2         Goldfinger 1964  125.0      Sean Connery
## 3 Diamonds are Forever 1971 120.0      Sean Connery
## 4           Moonraker 1979  210.3      Roger Moore
## 5 The Living Daylights 1987  191.2 Timothy Dalton
## 6           GoldenEye 1995  355.0 Pierce Brosnan
## 7       Casino Royale 2006  599.0   Daniel Craig
```

Tibble

- ▶ Similar to traditional data frame.
- ▶ tibble is the modern standard in R.

```
library(tibble)
myFavMovies2 <- tibble(title = c("Dr. No",
                                "Goldfinger",
                                "Diamonds are Forever",
                                "Moonraker",
                                "The Living Daylights",
                                "GoldenEye",
                                "Casino Royale"),
  year = c(1962, 1964, 1971, 1979,
           1987, 1995, 2006),
  box = c(59.5, 125, 120, 210.3,
          191.2, 355, 599),
  bondActor = c("Sean Connery",
                 "Sean Connery",
                 "Sean Connery",
                 "Roger Moore",
                 "Timothy Dalton",
                 "Pierce Brosnan",
                 "Daniel Craig"))

# Append an extra row at the end of the tibble
# Rewrite the original tibble object
myFavMovies2 <- add_row(myFavMovies2,
  title = "Spectre", year = 2015, box = 880.7,
  bondActor = "Daniel Craig")
```

Tibble

```
myFavMovies2
```

```
## # A tibble: 8 x 4
```

```
##   title                year    box bondActor  
##   <chr>                <dbl> <dbl> <chr>
```

```
## 1 Dr. No                1962  59.5 Sean Connery
```

```
## 2 Goldfinger            1964  125   Sean Connery
```

```
## 3 Diamonds are Forever  1971  120   Sean Connery
```

```
## 4 Moonraker             1979  210.  Roger Moore
```

```
## 5 The Living Daylights  1987  191.  Timothy Dalton
```

```
## 6 GoldenEye             1995  355   Pierce Brosnan
```

```
## 7 Casino Royale         2006  599   Daniel Craig
```

```
## 8 Spectre               2015  881.  Daniel Craig
```

Subsetting a Tibble

```
# Get one column by name  
myFavMovies2[["title"]]  
myFavMovies2$title  
# Get a range of columns by position ID  
myFavMovies2[, 1:2]  
myFavMovies2[1:2]  
# Get rows 1 to 3  
myFavMovies2[1:3, ]  
# Get the "year" variable of row 1-3  
myFavMovies2[1:3, "year"]  
# Get the "title" and "year" variables of row 4-7  
myFavMovies2[4:7, c("title", "year")]
```

Functions

- Functions are vectorised.

```
is.odd <- function(x) {  
  # The modulo operator %% returns the remainder  
  # If a number divide by 2 gives remainder 1, then it is an odd number  
  remainder <- x %% 2  
  equalToOne <- remainder == 1  
  return(equalToOne)  
}  
# Execute the function with one input  
is.odd(5)  
  
## [1] TRUE  
  
# Execute the function with an integer vector  
is.odd(1:10)  
  
## [1] TRUE FALSE TRUE FALSE TRUE FALSE TRUE FALSE TRUE FALSE  
  
# Define another function  
is.even <- function(x) {  
  !is.odd(x)  
}  
# Return true for even numbers  
is.even(1:10)  
  
## [1] FALSE TRUE FALSE TRUE FALSE TRUE FALSE TRUE FALSE TRUE
```


If-Else

```
library(lubridate)
# Find out what day is today
myWeekday <- weekdays(today())
# Check whether today is Saturday or Sunday
if (myWeekday %in% c("Saturday", "Sunday")) {
  myGreeting <- "Have a nice weekend"
} else {
  myGreeting <- "Go back to work"
}
# Prints the message
myGreeting
```

If-Else

- Multiple conditions are allowed.

```
library(lubridate)
myWeekday <- weekdays(today())
# Checks multiple conditions
if (myWeekday %in% c("Saturday","Sunday")) {
  myGreeting <- "Have a nice weekend"
} else if(myWeekday == "Friday"){
  myGreeting <- "It's Friday!"
} else if(myWeekday == "Monday"){
  myGreeting <- "Oh no..."
} else {
  myGreeting <- "Go back to work"
}
myGreeting
```

While

- ▶ Loops as long the condition stays TRUE.

```
myCounter <- 100
while (myCounter > 0) {
  myCounter <- myCounter - 5
  print(myCounter)
}
```

While

- ▶ Skip an iteration using `next`
- ▶ Early exit using `break`

```
myCounter <- 100
while (myCounter > 0) {
  myCounter <- myCounter - 5
  if (myCounter > 50) {
    # Skips all iterations if the counter value is greater than 50
    next
  }
  if (myCounter == 10) {
    # Early stop if the counter value matches 10
    break
  }
  print(myCounter)
}
```

```
## [1] 50
## [1] 45
## [1] 40
## [1] 35
## [1] 30
## [1] 25
## [1] 20
## [1] 15
```

For

```
myResult <- 0
for (i in 1:100) {
  myResult <- myResult + i ^ 2
}
myResult

## [1] 338350
```

Apply

- ▶ `apply()` takes a 'grid' input: data frame, tibble, or matrix
- ▶ Margin 1 = apply over rows
- ▶ Margin 2 = apply over columns

```
# The second argument 1 indicates iterate over rows
myMessages1 <- apply(myFavMovies2, 1, function(row){
  sprintf("%s was released in %s.", row["title"], row["year"])
})
myMessages1

## [1] "Dr. No was released in 1962."
## [2] "Goldfinger was released in 1964."
## [3] "Diamonds are Forever was released in 1971."
## [4] "Moonraker was released in 1979."
## [5] "The Living Daylights was released in 1987."
## [6] "GoldenEye was released in 1995."
## [7] "Casino Royale was released in 2006."
## [8] "Spectre was released in 2015."
```

lapply

- ▶ `lapply()` always returns a list

```
# The second argument 1 indicates iterate over rows
myMessages2 <- lapply(myFavMovies2$title,
                      function(x){ sprintf("%s is a great movie!", x) })
# Checks the data type of the result
typeof(myMessages2)

## [1] "list"

# Check whether it is a list
is.list(myMessages2)

## [1] TRUE

# Select the 6th element of the list
myMessages2[[6]]

## [1] "GoldenEye is a great movie!"
```

apply

- ▶ `sapply()` always returns a vector

```
myMessages3 <- sapply(myFavMovies2$title,  
                      function(x){ sprintf("%s is a great movie!", x) })  
# Check whether it is a list  
is.list(myMessages3)
```

```
## [1] FALSE
```

```
myMessages3
```

```
##                               Dr. No  
##      "Dr. No is a great movie!"  
##                               Goldfinger  
##      "Goldfinger is a great movie!"  
##                               Diamonds are Forever  
##      "Diamonds are Forever is a great movie!"  
##                               Moonraker  
##      "Moonraker is a great movie!"  
##                               The Living Daylights  
##      "The Living Daylights is a great movie!"  
##                               GoldenEye  
##      "GoldenEye is a great movie!"  
##                               Casino Royale  
##      "Casino Royale is a great movie!"  
##                               Spectre  
##      "Spectre is a great movie!"
```


Looping

- ▶ Looping can be slow.
- ▶ Always try to vectorise your code.

```
# Vectorised operation is fast
system.time({
  myResult <- 1:10000 * 2
})

##      user  system elapsed
##         0         0         0

# Looping is quite slow
system.time({
  myResult <- sapply(1:10000, function(x){ x * 2 })
})

##      user  system elapsed
##         0         0         0

# Appending to vector is much slower
system.time({
  myResult <- c()
  for(i in 1:10000){
    myResult <- c(myResult, i * 2)
  }
})

##      user  system elapsed
##    0.17    0.00    0.17
```

UK User Communities

- ▶ LondonR -  LONDONR
- ▶ ManchesterR -  MANCHESTERR
- ▶ CaRdiff -  CaRdiff
- ▶ SheffieldR
- ▶ EdinbR -  EdinbR
- ▶ CambR -  CambR
- ▶ NottinghamR
- ▶ BirminghamR -  BIRMINGHAMR
- ▶ Oxford RUG -  Oxford
- ▶ Bristol Data Scientists -  BRISTOL
DATA SCIENTISTS

International User Communities

- ▶ International R User Conference (useR!)



- ▶ Enterprise Application of the R Language (EARL)



- ▶ European R User Meeting (ERUM)




Data Transformation

Tidyverse

- ▶ A coherent system of packages for data manipulation, exploration and visualisation.
- ▶ Typical workflow of a project:

Program



Tidyverse: Packages

Import Reading datasets from various data sources.

- ▶ readr
- ▶ readxl
- ▶ haven
- ▶ httr
- ▶ rvest
- ▶ xml2

Tidy Clean up datasets.

- ▶ tidble
- ▶ tidyr

Transform Aggregate, change variable format and derive new variables.

- ▶ dplyr
- ▶ forcats
- ▶ hms
- ▶ lubridate
- ▶ stringr

Visualise Creating charts using the Grammar of Graphics.

- ▶ ggplot2

Model Train and test statistical models.

- ▶ broom
- ▶ modelr

Program Coding in pipeline-style.

- ▶ magrittr
- ▶ purrr

Data Transformation

- ▶ Subset the observations by condition - `filter()`
- ▶ Reorder the observations - `arrange()`
- ▶ Pick variables by name - `select()`
- ▶ Compute new variable as a function of existing variables - `mutate()`
- ▶ Aggregate many values into one - `summarise()` or `summarize()`
- ▶ All of the above functions can be used with `group_by()`
- ▶ Syntax: `function(data, actions_to_take)`

Loading Dataset

```
# Load the package
library(nycflights13)
# View the flights dataset interactively
View(flights)
# Print out the dataset on the console
flights

## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     1     544           545        -1    1004
## 5  2013     1     1     554           600        -6     812
## 6  2013     1     1     554           558        -4     740
## 7  2013     1     1     555           600        -5     913
## 8  2013     1     1     557           600        -3     709
## 9  2013     1     1     557           600        -3     838
## 10 2013     1     1     558           600        -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```


dplyr Query


- Finding average arrival delay time for each route and sort it.
(i.e. which route are delayed the most?)

```
library(dplyr)
flights %>%
  group_by(tailnum) %>%
  summarise(delay = mean(arr_delay),
            n = n()) %>%
  arrange(desc(delay)) %>%
  filter(n > 100)
```

```
## # A tibble: 1,201 x 3
##   tailnum delay    n
##   <chr>   <dbl> <int>
## 1 N826UA   19.6   120
## 2 N929DL   17.6   108
## 3 N431UA   16.3   104
## 4 N342NW   16.0   111
## 5 N567UA   15.0   105
## 6 N850UA   14.9   130
## 7 N179JB   14.3   311
## 8 N803UA   13.2   129
## 9 N561UA   12.8   103
## 10 N11206  12.7   111
## # ... with 1,191 more rows
```

Filtering Observations

- ▶ Pick a subset of observations based on their values
- ▶ Use logical operators: $>$, $>=$, $<$, $<=$, $!=$, $==$
- ▶ More complex combinations of logical operators:



- ▶ Only returns observations where the condition is TRUE. All FALSE and NA values are excluded.
- ▶ Use the function `is.na()` to check if the value is missing.

Examples - filter()

```
# Select all flights on January 1st  
# Assigning them to a new variable jan1  
jan1 <- filter(flights, month == 1, day == 1)  
jan1  
# Select all flights from November or December  
filter(flights, month == 11 | month == 12)  
# A useful shorthand for this problem is x %in% y  
# It selects every row where x is one of the values in y  
filter(flights, month %in% c(11, 12))  
# Let's select only the flights that weren't delayed  
# (on arrival or departure) by more than two hours:  
filter(flights, arr_delay <= 120, dep_delay <= 120)
```

Arranging Observations

- ▶ Changes the order of observations in a dataset.
- ▶ Sorts the observations by a set of variables in ascending order
 - ▶ Use `desc()` to sort in descending order
 - ▶ If more than one variable is supplied, each additional variable will break ties in the values of the preceding variable.
 - ▶ NA values are placed at the end

Examples - arrange()

```
# Arrange flights by year, then month, then day  
arrange(flights, year, month, day)  
# Use desc() to reorder by a column in descending order  
arrange(flights, desc(arr_delay))  
# Missing values are always sorted at the end:  
df <- tibble(x = c(5, 2, NA))  
arrange(df, x)  
arrange(df, desc(x))
```

Selecting Variables

- ▶ Returns specific variables in the dataset, dropping the others.
- ▶ Comes with a number of helper functions you can use within `select()` to pick out variables based on their names:
 1. `starts_with("abc")` matches variable names that begin with "abc".
 2. `ends_with("xyz")` matches variable names that end with "xyz".
 3. `contains("ijk")` matches variable names that contain "ijk".
 4. `num_range("x", 1:3)` matches `x1`, `x2`, and `x3`.
- ▶ Can be used to rename variables, but better to use the `rename()` function
- ▶ To move several variables to the start of a data frame use `select()` in conjunction with `everything()`

Examples - select()

```
# Select columns by name
select(flights, year, month, day)
# Select all columns between 'year' and 'day' (inclusive)
select(flights, year:day)
# Select all columns except those from year to day (inclusive)
select(flights, -(year:day))
# Rename a variable using rename()
rename(flights, tail_num = talinum)
# Reorder columns using the everything() helper
select(flights, time_hour, air_time, everything())
```

Compute New Variables

- ▶ Create new variables which are functions of existing ones
- ▶ New variables are always added at the end of the dataset
- ▶ To keep only the newly-computed variables and remove the old ones, use `transmute()`.
- ▶ Many functions can be used with `mutate()` to create new variables:

Arithmetic operators `+`, `-`, `*`, `/`, `^`

Modular arithmetic `%/%` for integer division and `%%` for modulo

Logs Very useful for data ranging across multiple orders of magnitude

Comparison `<`, `<=`, `>`, `>=`, `!=`, `==`

Ranking There are several of these, the most common one is `min_rank()`

Examples - summarise()

```
mySummary <- flights %>%  
  group_by(dest) %>%  
  summarise(  
    count = n(),  
    dist = mean(distance, na.rm = TRUE),  
    delay = mean(arr_delay, na.rm = TRUE)) %>%  
  filter(count > 20)  
mySummary
```

```
## # A tibble: 97 x 4  
##   dest count dist delay  
##   <chr> <int> <dbl> <dbl>  
## 1 ABQ    254 1826  4.38  
## 2 ACK    265  199  4.85  
## 3 ALB    439  143 14.4  
## 4 ATL  17215  757. 11.3  
## 5 AUS  2439 1514.  6.02  
## 6 AVL    275  584.  8.00  
## 7 BDL    443  116  7.05  
## 8 BGR    375  378  8.03  
## 9 BHM    297  866. 16.9  
## 10 BNA  6333  758. 11.8  
## # ... with 87 more rows
```

Grouped Summaries

- ▶ `summarise()` and `summarize()` aggregates a set of values into one.
- ▶ Commonly used with `group_by()` to analyse properties of individual groups.
- ▶ Examples of summary functions:
 - Measures of location Arithmetic average `mean()` and median `median()`.
 - Measures of spread Standard deviation `sd()` and interquartile range `IQR()`.
 - Measures of rank Minimum value `min()`, maximum value `max()` as well as the quantiles `quantile()`
 - Measures of position `first()` and `last()`
- ▶ Use the pipe operator `%>%` to combining several operations

Examples - mutate()

```
# Select several columns only
flights_sml <- select(flights,
                      year:day,
                      ends_with("delay"),
                      distance,
                      air_time)

# Use these the smaller data frame derive new columns
mutate(flights_sml,
       gain = arr_delay - dep_delay,
       speed = distance / air_time * 60)

## # A tibble: 336,776 x 9
##   year month   day dep_delay arr_delay distance air_time   gain speed
##   <int> <int> <int>     <dbl>     <dbl>     <dbl>   <dbl> <dbl> <dbl>
## 1  2013     1     1         2         11     1400     227      9  370.
## 2  2013     1     1         4         20     1416     227     16  374.
## 3  2013     1     1         2         33     1089     160     31  408.
## 4  2013     1     1        -1        -18     1576     183    -17  517.
## 5  2013     1     1        -6        -25      762     116    -19  394.
## 6  2013     1     1        -4         12      719     150     16  288.
## 7  2013     1     1        -5         19     1065     158     24  404.
## 8  2013     1     1        -3        -14      229      53    -11  259.
## 9  2013     1     1        -3         -8      944     140     -5  405.
## 10 2013     1     1        -2          8      733     138     10  319.
## # ... with 336,766 more rows
```

Regression Models

Simple Regression

Univariate linear regression model

$$\hat{y}_i = \beta_0 + \beta_1 x_i$$

- ▶ Analogous to a straight line $y = mx + c$
- ▶ Can be chained with M dependent variables (Multivariate)

$$\hat{y}_i = \beta_0 + \sum_{m=1}^M \beta_m x_{m,i}$$

- ▶ Residual term $\epsilon_i = y_i - \hat{y}_i$ assumed to be Gaussian

$$\epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

- ▶ Also known as ordinary least squared (OLS) regression

Linear Regression in R

```
# Build a univariate linear model
# These two lines are equivalent
myModel1 <- lm(mpg ~ wt, mtcars)
myModel1 <- lm(formula = mpg ~ wt, data = mtcars)
# Read the model summary
summary(myModel1)

##
## Call:
## lm(formula = mpg ~ wt, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.5432 -2.3647 -0.1252  1.4096  6.8727
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   37.2851     1.8776  19.858 < 2e-16 ***
## wt           -5.3445     0.5591  -9.559 1.29e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.046 on 30 degrees of freedom
## Multiple R-squared:  0.7528, Adjusted R-squared:  0.7446
## F-statistic: 91.38 on 1 and 30 DF,  p-value: 1.294e-10
```

More Linear Regression Models

Multivariate linear model Additional independent variables can be chained using the + symbol. Categorical variables can be encoded as dummy on-the-fly using `factor()`.

```
myModel2 <- lm(mpg ~ wt + hp + qsec + factor(am), mtcars)
```

Polynomial term Model can become more flexible when an independent variable is converted into polynomial terms. Use the `poly()` function.

```
myModel3 <- lm(mpg ~ wt + qsec + factor(am) +  
               poly(hp, 3), mtcars)
```

Interaction term Two variables can be combined to create synergy effect. The * symbol is used to combine variables together.

```
myModel4 <- lm(mpg ~ wt * hp + qsec + factor(am), mtcars)
```

Regression Diagnostics





Residuals vs Fitted Checks for non-linear relationship. Look for a near horizontal line.

Normal Quantile-Quantile It aligns model residuals against a theoretical normal distribution. If the residuals spread along a straight diagonal line on the Q-Q plot, it suggests that the residuals are normally distributed.

Scale-Location Checks for homoscedasticity and heteroscedasticity. It is homoscedastic if observations scatter without any observable pattern.

Residual vs Leverage (Cook's Distance) Identifies observations having strong influence to the model.

Diagnostics Plots




Overfitting

- ▶ Flexible models are prone to overfitting.
- ▶ Overfitting makes the model less generalisable.
- ▶ Solution
 - ▶ Use less flexible methods.
 - ▶ Impose regularisation.


Overfitting: Visual Explanation

$$\hat{y} = \beta_0 + \sum_{j=1}^8 \beta_{wtj} x_{wt}^j + \sum_{k=1}^5 \beta_{hp_k} x_{hp}^k$$



Poisson Distribution

- ▶ Count of distinct events are drawn from Poisson distribution.
 - ▶ Always positive.
 - ▶ In most cases they are integers.
- ▶ e.g. Number of people in a room, number of flights delayed per day... etc.




Testing for Poisson Distribution

- ▶ Chi-square goodness of fit test.
- ▶ Fits the data against theoretical Poisson distribution.
- ▶ Look for statistical significance.

```
# Performs the Chi-squared goodness-of-fit test.  
# It checks whether the variable is drawn from a Poisson distribution.  
library(vcd)  
gf <- goodfit(mtcars$carb, type= "poisson", method= "ML")  
# Checks the statistical p-value of the goodness-of-fit test.  
# If  $p \leq 0.05$  then it is safe to say that the variable is Poisson.  
summary(gf)  
  
##  
## Goodness-of-fit test for poisson distribution  
##  
## X^2 df P(> X^2)  
## Likelihood Ratio 20.53973 4 0.0003906369
```

Goodness of Fit Plot

Plots the observed frequency vs theoretical Poisson distribution.
The hanging bars should fill the space if it is perfectly Poisson.
`plot(gf)`



Poisson Regression

```
# Build a Poisson model to predict the number of carburetors in a car.
myPoissonModel <- glm(carb ~ hp + wt + factor(am),
                      family="poisson",
                      data=mtcars)


# Read the model summary
summary(myPoissonModel)

##
## Call:
## glm(formula = carb ~ hp + wt + factor(am), family = "poisson",
##      data = mtcars)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.91420  -0.48423  -0.07246   0.19252   1.26155
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.418081   0.604211  -0.692   0.4890
## hp           0.004316   0.001880   2.296   0.0217 *
## wt           0.179583   0.191352   0.938   0.3480
## factor(am)1  0.393750   0.324978   1.212   0.2257
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 27.043  on 31  degrees of freedom
## Residual deviance: 10.798  on 28  degrees of freedom
## AIC: 108.16
##
## Number of Fisher Scoring iterations: 4
```

Binomial Distribution

- ▶ There are only two possible outcomes $\{Y, \neg Y\}$.
 - ▶ Toss a coin (Head or tail)
 - ▶ Taking an examination (Pass or fail)
 - ▶ Selling a product (Sold or unsold)
- ▶ Likelihood of event Y and $\neg Y$ expressed as probability $P(Y) + P(\neg Y) = 1$.
- ▶ Logistic function squeezes real value range X into $(0, 1)$ to express probability.

$$P(Y) = \frac{1}{1 + e^{-X}}$$



Logistic Regression

- Equation for logistic regression

$$P(Y) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_M x_M)}}$$

- Coefficients $\beta_1, \beta_2, \beta_3, \dots, \beta_M$ can be converted into odds ratios $OR(x_1), OR(x_2), OR(x_3), \dots, OR(x_M)$.

$$OR(x_1) = \frac{odds(x_1 + 1)}{odds(x_1)} = \frac{e^{\beta_0 + \beta_1(x_1 + 1) + \beta_2 x_2 + \dots + \beta_M x_M}}{e^{\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_M x_M}} = e^{\beta_1}$$

- $OR(x_1)$ Represents the change in probability when x_1 increases by 1 unit.

Training a Logistic Regression Model

- Build a logistic regression model to predict the dependent variable am. (1=manual; 0=auto)

```
myLogisticModel <- glm(am ~ mpg + hp + disp,  
                        family="binomial",  
                        data=mtcars)  
summary(myLogisticModel)
```

Calculate the odds ratios for this model.


```
# Calculate the odds-ratios by taking the exponential of the coefficients  
# You may also calculate the 95% confidence interval of the odds-ratio  
exp(cbind(oddsratio = myLogisticModel$coefficients,  
          confint(myLogisticModel)))
```

| ## | oddsratio | 2.5 % | 97.5 % |
|----------------|--------------|--------------|-------------|
| ## (Intercept) | 2.066683e-15 | 9.104693e-49 | 0.02539238 |
| ## mpg | 3.614582e+00 | 1.198283e+00 | 61.43800377 |
| ## hp | 1.161095e+00 | 1.049568e+00 | 1.50974362 |
| ## disp | 9.366496e-01 | 8.197068e-01 | 0.98685201 |

Tree-based Methods

Recursive Partitioning

- ▶ Cut off point is denoted as s .
- ▶ Divides data into regions (leaves) $\mathcal{R}_1, \mathcal{R}_2, \mathcal{R}_3, \dots$ recursively.
- ▶ Works with real values as well as categorical variables.
- ▶ Large tree risks overfitting
 - ▶ Removes weaker leaves.
 - ▶ Regularisation.




Decision Tree

- Trees can be trained with a formula and optional control parameters.

```
# Load the rpart package for recursive partitioning  
library(rpart)  
# Build a decision tree to predict mpg  
myTree <- rpart(formula = mpg ~ wt + hp +  
                 factor(carb) +  
                 factor(am),  
                 data = mtcars,  
                 control = rpart.control(minsplit=5))  
# Read the detailed summary of the tree  
summary(myTree)
```

Decision Tree: Visualisation

```
# Load the rpart.plot package for tree visualisation  
library(rpart.plot)  
rpart.plot(myTree)
```




Tree Pruning

```
printcp(myTree)

##
## Regression tree:
## rpart(formula = mpg ~ wt + hp + factor(carb) + factor(am), data = mtcars,
##       control = rpart.control(minsplit = 5))
##
## Variables actually used in tree construction:
## [1] factor(carb) hp          wt
##
## Root node error: 1126/32 = 35.189
##
## n= 32
##
##      CP nsplit rel error  xerror   xstd
## 1 0.652661    0  1.000000 1.08570 0.252677
## 2 0.178618    1  0.347339 0.69922 0.172943
## 3 0.046269    2  0.168721 0.43098 0.098563
## 4 0.026109    3  0.122453 0.34614 0.095951
## 5 0.022593    4  0.096343 0.36810 0.097621
## 6 0.011989    5  0.073751 0.36754 0.087290
## 7 0.010000    6  0.061762 0.35358 0.081453

plotcp(myTree)
```




Random Forest

- ▶ Consists of many decision trees
 - ▶ Randomly selected variables will be used in each split
 - ▶ Usually no need to prune them (all trees are allowed to grow big)
- ▶ M trees in a random forest produces M predictions
 - ▶ Final prediction is calculated as mean value for regression problem
 - ▶ Classification problem will use most the common label (majority voting)

Training a Random Forest

```
library(randomForest)
library(dplyr)
# Build a random forest with 1000 trees
# Each tree has 2 randomly selected variables
# You can change the parameters
myForest <- randomForest(mpg ~ wt + hp + carb + am,
                          ntree = 1000,
                          mtry = 2,
                          data = mtcars %>% mutate(carb = factor(carb),
                                                    am = factor(am)))



# Plot the error as the forest expands
plot(myForest)
```



Neural Networks

Artificial Neurons

- ▶ Inspired by neurons in biological brain.
- ▶ McCulloch and Pitts described neuron as a logical process.
 - ▶ Neuron takes several inputs $\{x_1, x_2, x_3, \dots, x_M\}$
 - ▶ Fires (activates) if the combined weighted input $\sum_{m=1}^M w_m x_m$ exceeds threshold.
 - ▶ Output can either be fire 1 or not fire 0.
- ▶ Rosenblatt's Mark I Perceptron



Modern Neural Networks

- ▶ Neural nets are based on non-linear processing power.
- ▶ Trained via gradient descent optimisers (backpropagation).
 - ▶ Network initialise randomly.
 - ▶ Requires differentiable loss function.
 - ▶ Requires strong gradient in order to improve.
 - ▶ Model weights improve iteratively.
 - ▶ Converge at local minimum.
- ▶ Neurons can be stacked as layers.
 - ▶ Can either be shallow (1 layer) or deep (many layers).
 - ▶ State-of-the art neural networks have highly bespoke topologies.


Activation

- ▶ Weighted inputs are combined linearly.

$$X = \sum_{m=1}^M w_m x_m$$

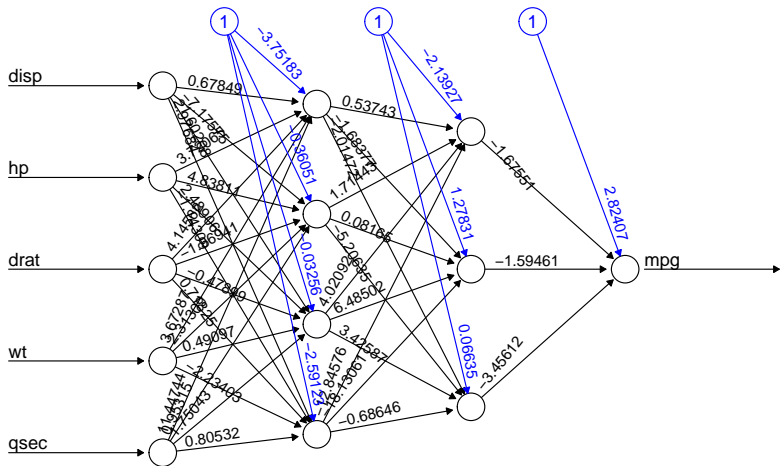
- ▶ Non-linear activation functions

- ▶ Sigmoid
- ▶ Hyperbolic tangent
- ▶ etc...



Multilayer Perceptron: Topology

- ▶ Layers are fully-interconnected.
- ▶ Usually having two or more layers.



Error: 0.045594 Steps: 2263

Time Series Analysis

Time Series Data

- ▶ Observations repeatedly taken at regular interval.
- ▶ Explore variable relationship across temporal space.




Auto-correlation Function (ACF) Measures the correlation of a single variable along the temporal dimension between x_t and x_{t+h} . It shows the correlation of the variable over different lag periods. For most time series variables, correlation is usually strong at lag $h = 1$ and it gradually diminishes as lag period increases. Cyclic pattern in the correlogram suggests possible seasonality which you can analyse further.

Partial Auto-correlation Function (PACF) Also measures the correlation between different lag periods, but it controls the correlation across the temporal dimension so that only the contribution of an individual lag is reflected.

Cross Correlation Function (CCF) Analyses the temporal correlation between two variables.

ACF and PACF Correlograms


myTrainingSet\$total_demand



Decomposition

- ▶ Time series can be decomposed into:
 - ▶ Seasonal component S_t
 - ▶ Trend component T_t
 - ▶ Residual component ϵ_t
- ▶ Additive time series is expressed as $X_t = S_t + T_t + \epsilon_t$
- ▶ Multiplicative time series is expressed as $X_t = S_t \times T_t \times \epsilon_t$

Decomposition of additive time series



Time Series Linear Regression Model

- ▶ Decomposed components can be used as independent variable in linear regression:

$$X_t = \beta_0 + \beta_{trend} T_t + \beta_{seasonal} S_t + \sum_{m=1}^M (\beta_m x_{mt}) + \epsilon_t$$

- ▶ Components can also form interaction terms with other independent variables.

Auto-regressive Moving Average Model (ARMA)

- ▶ $ARMA(p, q)$ model


$$\underbrace{X_t}_{\text{Observation}} = \underbrace{\beta_0}_{\text{intercept}} + \underbrace{\sum_{i=1}^p (\phi_i X_{t-i})}_{\text{AR}(p)} + \underbrace{\sum_{i=1}^q (\theta_i \epsilon_{t-i})}_{\text{MA}(q)} + \underbrace{\epsilon_t}_{\text{residual}}$$

- ▶ $ARIMA(p, d, q)$ model

- ▶ Auto-regressive Integrative Moving Average
- ▶ $I(d)$ d^{th} order integration can be added.
 - ▶ Integration refers to the difference from previous time step
 - ▶ First order differencing $I(1)$: $X'_t = X_t - X_{t-1}$
 - ▶ To satisfy **stationarity** requirement.

Stationarity


- ▶ Equal properties (mean and variance) across time.
 - ▶ Only one below is a stationary time series.
 - ▶ Which one is it?



ARIMA with Seasonality (SARIMA)

- ▶ $ARIMA(p, d, q)(P, D, Q)_m$
 - ▶ All parameter values can be automatically identified.
 - ▶ Simple models are always preferred
 - ▶ Intend to keep $p + q + P + Q$ small.

Forecasts from Regression with $ARIMA(2,0,0)(1,1,1)[7]$ errors



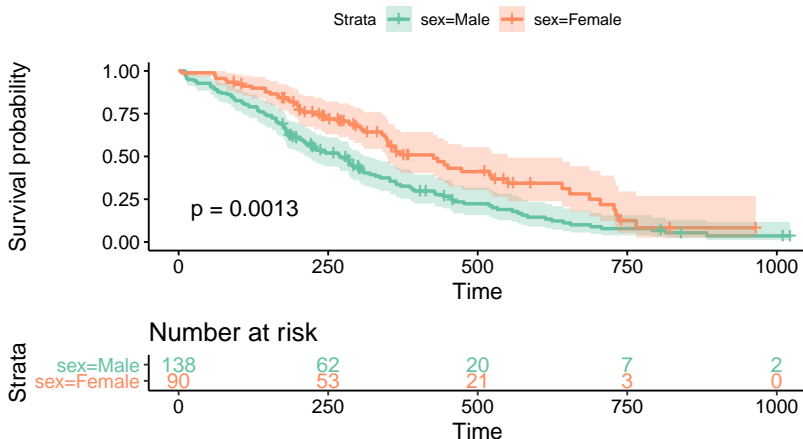
Survival Analysis

Survival Analysis

- ▶ Event occurring at irregular intervals.
 - ▶ e.g. Patients gets sick, machine failing... etc
- ▶ Also known as **time-to-event analysis** or **event history analysis**.

Kaplan-Meier Estimator

- ▶ It is used to measure how many subjects survives in a clinical trial since treatment began.
- ▶ Categorical variables only.



Cox Proportional Hazard Model

- ▶ It is a regression method which can take into account categorical and numeric variables.
- ▶ Assumes effects are time-independent (proportional hazard assumption).
- ▶ Hazard function h_t is defined as:

$$h_t = h_{0,t} \times e^{\beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_M x_M}$$





```
# Build a Cox model with predictor variables
myCoxModel1 <- coxph(Surv(time, status==2) ~ factor(sex) + age +
                    ph.ecog + ph.karno +
                    pat.karno +
                    meal.cal + wt.loss, data = lung)

# Read the model summary
summary(myCoxModel1)
```

Unsupervised Learning

K-means Clustering

- ▶ Works with unlabelled data.
- ▶ Arrange objects into groups (clusters)
- ▶ Objective interpretation of results as K is arbitrarily selected.



Agglomerative Hierarchical Clustering

- ▶ N objects can form maximum N clusters, each having 1 member object.
- ▶ Identify distance between closest cluster pair.
- ▶ Merge them.
- ▶ Repeat until there are no more clusters left.

Agglomerative Hierarchical Clustering: Example

